1. **Problem Statement:** Predicting insurance charges based on input variables such as age, sex, BMI, number of children, and smoking status, with insurance charges as the output variable.

2. **Basic Info about Dataset:**

   - Total number of rows and columns not mentioned.

   Input Variables:

   - Age (continuous)
   - Sex (categorical)
   - BMI (continuous)
   - Children (continuous)
   - Smoker (categorical)

   Output Variable:

   - Charges (continuous)

3. **Pre-processing Method:** One-hot encoding using the **get_dummies** method in Python for converting categorical variables into numerical form.

4. **Model Development with R-squared- SVM model:**

   - Support Vector Machine (SVM)
     - Various hyperparameters (C values) tested for different kernels: Linear, RBF, Polynomial, and Sigmoid.
     - R-squared values documented for each model.

S.No	Hyper Parameter	Linear	RBF (Non-Linear)	Polynomial	Sigmoid
1	C10	0.46593	0.0331	0.03906	-7.33
2	C100	0.63124	0.3137	0.6156	-0.01756
3	C500	0.7644	0.6619	0.8234	0.1254
4	C1000	0.7671	0.8114	0.852	0.287
5	C2000	0.7478	0.8583	0.85623	-0.5878
6	C5000	0.74536	0.8784	0.8538	-7.33

5. **Research Values (R-squared scores):**

6. The RBF kernel consistently improves R-squared value with increasing C values, peaking at C5000.

7. **Final Model:** The chosen final model utilizes the RBF kernel with C=5000, achieving the highest R-squared value of 0.8784. This model was selected for its superior performance compared to other hyperparameters and kernels, ensuring a better fit to the data and potentially higher prediction accuracy.

Model Development with R-squared - Decision Tree Model:

Decision Tree (DT)

R- Research Values (R-squared scores):

- The CRITERION, MAX FEATURES, and SPLITTER consistently improve R-squared value with increasing R Squared values.

The model configuration with max features set to 'auto' and splitter set to 'random' achieved an R value of 0.74420, demonstrating a notable level of performance.

Final Model:

The chosen model configuration with the criterion set to 'Mae', max features set to 'auto', and splitter set to 'random' achieved an R value of 0.74420, demonstrating a strong level of performance. This configuration was selected for its notable performance, ensuring a reliable fit to the data and potentially higher prediction accuracy.

S.No	CRITERION	MAX FEATURES	SPLITTER	R VALUE
1	Mse	auto	best	0.7083
2	Mse	auto	random	0.6782
3	Mse	sqrt	best	0.6488
4	Mse	sqrt	random	0.6307
5	Mse	log2	best	0.7075
6	Mse	log2	random	0.5103
7	Mae	auto	best	0.6787
8	Mae	auto	random	0.74420
9	Mae	sqrt	best	0.7050
10	Mae	sqrt	random	0.6272
11	Mae	log2	best	0.6398
12	Mae	log2	random	0.6973
13	frideman_mse	auto	best	0.68097
14	frideman_mse	auto	random	0.7139
15	frideman_mse	sqrt	best	0.70823
16	frideman_mse	sqrt	random	0.7041
17	frideman_mse	log2	best	0.7486
18	frideman_mse	log2	random	0.6848

Interpreting the Hyperparameter Tuning Results for Random Forest:

The provided table summarizes the performance of various Random Forest models with different hyperparameters. The key metric for performance is the R-squared value, with higher values indicating better model performance.

Best Hyperparameters:

From the table, we identify the configurations with the highest R-squared values:

S.No	n_estimators	max_features	max_depth	min_samples_split	min_samples_leaf	bootstrap	criterion	R-squared
1	100	auto	None	2	1	TRUE	mse	0.82905
2	100	sqrt	None	2	1	TRUE	mse	0.8669
3	100	log2	None	2	1	TRUE	mse	0.8409
4	200	auto	20	2	1	TRUE	mse	0.8369
5	200	sqrt	20	2	1	TRUE	mse	0.8256
6	200	log2	20	2	1	TRUE	mse	0.7466
7	200	auto	None	10	4	TRUE	mse	0.8329
8	200	sqrt	None	10	4	TRUE	mse	0.8274
9	200	log2	None	10	4	TRUE	mse	0.8141
10	500	auto	20	2	1	FALSE	mae	0.8315
11	500	sqrt	20	2	1	FALSE	mae	0.8285
12	500	log2	20	2	1	FALSE	mae	0.8252
13	500	auto	None	10	4	FALSE	mae	0.8104
14	500	sqrt	None	10	4	FALSE	mae	0.8138
15	500	log2	None	10	4	FALSE	mae	0.801

- **Model 2:**
  - n_estimators: 100
  - max_features: sqrt
  - max_depth: None
  - min_samples_split: 2
  - min_samples_leaf: 1
  - bootstrap: TRUE
  - criterion: mse
  - R-squared: 0.8669

Final Model Selection:

Based on the highest R-squared value, Model 2 is selected as the final model for deployment.

Summary of the Selected Model:

- Number of Estimators: 100
- Max Features: sqrt (square root of the total number of features)
- Max Depth: None
- Min Samples Split: 2
- Min Samples Leaf: 1
- Bootstrap: TRUE
- Criterion: mse (Mean Squared Error)
- R-squared: 0.8669

This model configuration is chosen due to its superior performance in terms of the R-squared value, indicating a better fit to the data and hence, more reliable predictions.

MultiLinear Regression Model:

In this regression model, identified the following coefficients:
• The first coefficient is approximately 257.80.
• The second coefficient is approximately 321.06.
• The third coefficient is approximately 469.58.
• The fourth coefficient is approximately -41.75.
• The fifth coefficient is approximately 23418.67.

Coefficients: These numbers represent how much each independent variable affects the predicted outcome. For example, a higher coefficient means that variable has a stronger impact.

Intercept: This value (-12057.24) is what the model predicts when all independent variables are zero. It's like the starting point of our predictions.

R-squared: This is a measure of how well our model fits the data. The higher the R-squared value (0.7895 in this case), the better our model explains the variation in the predicted outcome. So, around 78.95% of the variability in the outcome can be explained by our model's variables.